

# CSC 577 HW1

Simon Swenson

1/16/2019

## 1 Introduction



Figure 1: The original tent image, a key point of focus for this homework assignment.

All homework problems (including grad problems) were completed. I had some MATLAB experience from Probabilistic Graphical Models last semester, so this was not that time consuming, but it was a bit tedious, as I needed a refresher, and learning the syntax and features of any language is always a bit tedious.

## 2 (#6)Documentation for the Code of HA1

Below is the output of "help ha1" which documents the functionality of the code used for this assignment:

— help for hw1 —

hw1 Performs all commands needed for hw1. The return value is the number of completed homework questions. First, hw1 reads in an image, displays it in a figure, then saves the image with the extension '.out.jpg'. Then, it displays statistics (min, max) for each channel in the image. Then, the image is converted to a grayscale. This grayscale image is displayed as a figure. After that, each separate channel is taken as a grayscale image and displayed. To further experiment with channels, a new 3-channel image is produced which permutes the channels of the original image. To experiment with indexing, the grayscale image from earlier is modified so that each fifth pixel along the rows and columns are turned white. This is also displayed in a figure. In addition to the statistics gathered earlier, a histogram is then generated for each channel and shown in a figure. After that, to experiment with function plotting, the sin and cosine function are

plotted from the domain  $-\pi$  to  $\pi$ . To experiment with matrix operations, an arbitrary matrix equation is solved by using `inv`, `linsolve`, and `'`. The topic of modifying certain pixel values in the grayscale image is then returned to, and a one-liner is used to perform a similar operation to the lattice operation that the nested loop performed earlier. A similar method is then used to set all pixel values over a certain threshold to black. Finally, the function explores PCA and transforms a set of data points based on their PCA.

### 3 (#7) Basic Image Data Structures

A cursory look at the image data after it has been loaded (by using the `"whos"` function) reveals that the image is stored as a three dimensional matrix/tensor with dimensions representing y coordinate, x coordinate, and channel (r, g, b), in that order. The size of the tensor is  $489 \times 728 \times 3$ . All elements in that matrix/tensor are of type `"uint8."` To hold onto that size information for future use, we can assign multiple variables at the same time, since `"size"` returns an n-tuple.

In addition to the dimensions of the image, other useful statistics can be gathered. For example, by using `"min"` and `"max,"` we find that the minimum and maximum values are as follows:

- Minimum of red = 0
- Minimum of green = 0
- Minimum of blue = 0
- Maximum of red = 251
- Maximum of green = 248
- Maximum of blue = 253

At first, I was surprised at the minimum findings, as images usually do not contain true black. However, this finding does not necessarily mean that there is a (0, 0, 0) pixel in the image, just that 0 appears for each channel in at least one pixel, not necessarily the same pixel for all three channels.

Finally, the image was converted to a grayscale image. This created a great point of comparison for the following task, which examined each individual channel in the image *as if it were a grayscale image*.

### 4 (#8) Image Channels

For this experiment, each channel of the image was converted to a grayscale image separately. This really illustrated the effect of each channel on the image, especially the tent. However, much of the image is desaturated, which made it a bit harder to draw findings from this experiment.




Figure 2: Each channel of the original image, represented separately as if they were grayscale. Left: red channel, center: green channel, right: blue channel.

The most striking difference between the three channels is the amount of red and green in the tent when compared to the amount of blue. This illustrates that yellow can be reproduced for human eyes by simply combining equal amounts of red and green light. As the yellow in the image is so saturated, very little blue is

present, but, as we can see in the desaturated rock cliffs and snow, as colors get closer to gray, each channel starts to have an equal "say" in the final value. One final note, which may seem a bit obvious, is that white is reflected by a high value for each of the three channels, and black is reflected by a low value for each of the three channels.

Another interesting experiment involving the color channels of the image is to simply permute the channels and display the resulting color image. The most important thing this reveals is that a large red value and a large blue value create purple, as can be seen in the purple tent.



Figure 3: A new image created by permuting the channels of the original tent image. In this image, green was moved to red, blue was moved to green, and red was moved to blue. This reveals that a high red and a high blue value produce purple.

## 5 (#9) Manipulating Matrices

MATLAB, as its name suggests, is an extremely useful tool for dealing with matrices. Thus, the various matrix operations are important to learn. For this experiment, we used a for loop to address every fifth pixel of the grayscale image, both horizontally and vertically, and turn it white. The results follow in two different formats.

The two display modes of the data have their strengths and weaknesses. To view the image as it actually is, "imshow" is the obvious choice. It displays 2 dimensional matrices as gray images and 3 dimensional tensors as color images. This corresponds very nicely with the output from "imread." However, the second display method could be fruitful for visualizing data which does not necessarily correspond to the wavelengths of light we can normally see. For example, a similar hue-based heatmap approach is used in infrared cameras to visualize radiant heat. In this case, however, I think the best visualization is the grayscale image, as the backing data is derived from wavelengths that we can actually see. In addition, the small white pixels are harder to see when the image is squished as it is in "imagesc."

## 6 (#10) Histograms

While the min and max values of each channel above are a good start for understanding the image, there are more robust statistics that we can use. Histograms give us a good overview of the *concentration* of various color values for each of the three channels.






Figure 4: The resulting lattice of white dots when shown using the "imshow" function.




Figure 5: The resulting lattice of dots when shown using the "imagesc" function.

The histogram is actually very reflective of the broad regions of the image. This image is very simple, and has three broad regions: cliffs, the tent, and snow. All three histograms show a very large hump near 50, reflecting the large area of cliffs in the image. The tent and snow are a little harder to tease apart. Comparing the red histogram to the blue histogram gives us some hints, however. The peak of the blue hump is at around 175. To me, this indicates that the peaks for the snow will be at 175 for the other two channels. However, the peaks for red and green are actually closer to 130-150. This is due to the tent. This illustrates one weakness of histograms: if two regions of the image have colors that are close in value to one another, you will not necessarily be able to distinguish those regions in the histogram.




Figure 6: Three histograms, each representing the concentration of values for a given channel. Left: red, center: green, right: blue.

## 7 (#11) Plotting

The following is a very simple plot to get used to how plotting is done in MATLAB.




Figure 7: The plots of sin (blue) and cosine (red) from  $-\pi$  to  $\pi$ .

## 8 (#12) Playing with Linear Algebra

A good grasp of linear algebra is necessary for success in this course. Luckily, MATLAB has good builtin functions for performing common linear algebra tasks. This experiment involved using those builtins to solve for  $x$  in a system of linear equations of the form  $Ax = y$ . The matrix can be represented as:

$$\begin{bmatrix} 3 & 4 & 1 \\ 2 & -1 & 2 \\ 1 & 1 & -1 \end{bmatrix} x = \begin{bmatrix} 9 \\ 8 \\ 0 \end{bmatrix}$$

To solve this equation,  $x$  must be isolated. Thus, the matrix  $A$  must be inverted. MATLAB has a number

of methods to do so: "inv," "linsolve," and the " " operator.

The answer using "inv" was:

$$x = \begin{bmatrix} \sim 1.9375 \\ \sim 0.2500 \\ \sim 2.1875 \end{bmatrix}$$

To verify this, we plug it into the original equation and perform matrix multiplication, which yields:

$$\begin{bmatrix} 3 & 4 & 1 \\ 2 & -1 & 2 \\ 1 & 1 & -1 \end{bmatrix} \begin{bmatrix} \sim 1.9375 \\ \sim 0.2500 \\ \sim 2.1875 \end{bmatrix} = \begin{bmatrix} 9 \\ 8 \\ 0 \end{bmatrix}$$

An interesting side-effect of using floating point arithmetic for this task is that answers are not usually exact. This would be possible with arbitrary precision arithmetic, but, often, computer hardware, especially GPUs, is much faster at performing floating point operations than operations on arbitrarily large numbers. Comparing the answer above with the answer provided by linsolve, we see that they are not exactly the same. In fact, the difference between them is:

$$1.0 \times 10^{-15} \begin{bmatrix} \sim 0.2220 \\ \sim -0.0278 \\ 0 \end{bmatrix}$$

This could be the result of doing the operation in two steps by first inverting  $A$ , then multiplying  $A^{-1}y$ , which could cause cascading error, whereas "linsolve" might be optimizing in some other way. However, the error value is negligible, in this case.

## 9 (#15) Matrix Manipulations Without Explicit Loops

Earlier, we explored using a nested for loop to index particular (x, y) coordinates of the image. However, this task can be done much more simply by using a conditional within the index expression. This will result in only indices which meet the condition being used. This is similar to a feature in python's pandas library, and the syntax is very similar. This MATLAB feature was used to produce the following two figures:

## 10 (#16) PCA

We briefly studied PCA in my undergraduate linear algebra, but that was a while ago, so this review was welcome. An example dataset to which PCA will be applied is shown below:

The corresponding covariance matrix is:

$$C = \begin{bmatrix} \sim 0.0211 & \sim 0.0356 \\ \sim 0.0356 & \sim 0.0608 \end{bmatrix}$$

To me, it seems like this data could be represented much more naturally by rotating it either -60° or 30° so that the data varies substantially over only one axis rather than both axes. We could let the new basis vectors be something like:

$$b_1 = \begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$$

$$b_2 = \begin{bmatrix} \frac{-1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$$

In this case, I would expect the covariance matrix to be diagonal and have a large value like 1 in the upper-left and a small value like 0.05 in the bottom-right. This is because the first dimension would span a much larger range than the second dimension.



Figure 8: Here, we use conditionals within the index expression to address every fifth x and y, then set its color to black.




Figure 9: After setting each fifth x and y to black, a similar expression was used to address pixels with a value over a certain threshold, then set the values of those pixels to black.

## 11 (#17) PCA II

It turns out that the new basis vectors are actually the eigenvectors of the covariance matrix. Thus, we can use the "eig" builtin function. Doing so results in a matrix in which each column is one of those eigenvectors:

$$E = \begin{bmatrix} \sim 0.0211 & \sim 0.0356 \\ \sim 0.0356 & \sim 0.0608 \end{bmatrix}$$

To verify that the eigenvectors are orthogonal, we can simply verify that their dot product is zero, which turns out to be true. Then, we can simply transform the data by multiplying the original data set with the eigenvector matrix on the right. This yields the following plot:




Figure 10: An example data set. Note that, it appears to have a very stark cluster, but that cluster is not orthogonal to the axes. PCA attempts to transform this dataset so that the principal components are orthogonal to the axes. In this case, it means that a vector pointing to the top-right will be the first principal component.




Figure 11: A new coordinate system, transformed to the basis of the eigenvectors. Note that the points are now essentially arranged orthogonally to the x axis.



The resulting covariance matrix is diagonal, as we had hypothesized:

$$C = \begin{bmatrix} \sim 0.0818 & 0 \\ 0 & \sim 0.0001 \end{bmatrix}$$

The values that I had anticipated above were much different. I should have paid closer attention to the data range, which was less than one. This would yield a standard deviation of less than one, and thus a covariance even smaller still. However, the difference in magnitude of the two non-zero values in the covariance matrix was correctly postulated. One value is large and one is very small. Strikingly, the sums of the diagonals of the two covariance matrices (before transformation and after transformation) are the same. To me, this is unexpected, because the other values in the original covariance matrix seem to disappear. I would have expected those values to somehow get incorporated into the diagonals of the new covariance matrix, but that is probably a lack of understanding of what covariance really means.