

CSC 577 HW2

Simon Swenson

1/23/2019

1 Introduction

The main goal of this assignment was to get more familiar with spectral samples, sensors, the resultant RGB values, and the interplay between those three variables. The main activity was exploring methods of solving linear equations with multiple solutions and applying those methods to solve for hidden sensor sensitivities. I found sifting through the large amount of plots and error values for the undergraduate section tedious, but the concepts introduced in the graduate section were really cool. All problems were completed.

2 (#1) The Setup and Visualization




Figure 1: Sensor responses (RGB values) for a randomly generated set of 1600 spectra, laid out in a 40 x 40 grid.

This portion of the assignment really just involved reading in the data, performing the standard spectra times sensor operation, and visualizing the resultant RGB values. The assignment mentioned that the sensors were biased toward blue, but I did not expect such a large bias! There are very few hues in here that one could consider *not* blue. I see a few desaturated blue grays and aquas, though.

3 (#2) Predicting Sensors from Perfect Output Data

Now that we have the data, we can start exploring different ways to manipulate it. First, we pretend that we do not have the original sensor sensitivity data, and we try to recover it from the sensor responses (RGB values). Recall that, with no error, we could solve the linear system of equations by simply multiplying the RGB on the left by the inverse of the spectra matrix. The pseudoinverse method is instead used, since we have more spectral samples than spectral wavelength buckets. This has the handy side-effect that the output is also the solution with the lowest mean squared error.

In the case of perfect output (no noise or clipping), we'll find that the error is extremely low and that the sensors can basically be fully recovered. A small amount of error is present when we round to the RGB values to the nearest integer, as I had done in an earlier version of my program, but, if we do not perform any rounding, the curves are basically perfect matches. (In the current version of my program, the values are rounded only just before they are to be converted into an image to maintain as much information as possible for as long as possible.) This is to be expected, since the pseudoinverse method is meant to minimize the least squared error, and if we haven't introduced any noise, then there won't be any error. If we visualize it as a one-dimensional linear regression, like in class, all the points will lie exactly on the line. In reality, since MATLAB is using some floating point type (either float or double) under the hood, some negligible error is introduced.




Figure 2: Actual sensors (blue) and predicted sensors (red) for the three channels: red (left), green (center), blue (right). Note that the match is so good that the predicted sensors completely cover the actual sensors. In an earlier version of the program, this was performed after rounding to ints, which introduced some noise and allowed both actual and predicted sensor curves to be visible.




Figure 3: The same curves by predicting after int rounding. Notice the increased error due to the decreased precision.

The error values were as follows (for the non-int-rounded version):

	Red	Green	Blue
Sensor RMSE	7.51e-11	1.033e-10	2.906e-10
Output RGB RMSE	3.78e-13	5.32e-13	1.442e-12

4 (#3) Adding Gaussian Noise

Perfect output like that of above is not very realistic, so we instead explore output with Gaussian noise. First, we randomly compute values from a normal distribution with mean 0 and standard deviation 10, then, we add those values to the computed RGB values above. Like before, we pretend that we don't have the original sensor sensitivities, so we use the pseudoinverse method to find the least-squared-error solution. The plots of the actual sensors and predicted sensors is below.




Figure 4: When we add Gaussian noise with standard deviation 10, the predicted sensors are much noisier as well. However, they still roughly follow the sensor curves.

And the accompanying errors:

	Red	Green	Blue
Sensor RMSE	1038.2	996.3	874.1
Output RGB RMSE	2.7434	2.7259	2.3788

However, we can also clip RGB to be within the range $[0, 255]$, and this is much closer to reality, as images normally store RGB as 8-bit unsigned integers. Before looking at the plots, I would guess that the spectra most likely to be clipped will be spectra with a high value at either the peak or the tails of the sensor sensitivity curves. Thus, I would expect the predicted sensor curves to exhibit more noise at the tails and peak after clipping. Let's take a look:




Figure 5: Here, we clip the output RGBs to be in range $[0, 255]$ to more closely resemble reality and visualize the output sensors.

And the accompanying errors:

	Red	Green	Blue
Sensor RMSE	1038.2	996.3	880.1
Output RGB RMSE	2.7434	2.7259	2.3895

It turns out the noise plots and the errors are very similar versus the unclipped data. In this case, this is probably indicative that there are very few color values close to the extremes. Remember that we need both an extreme value and Gaussian noise in the right direction to push it over the boundary.

5 (#4) Exploring Higher Noise Levels

The noise does add a significant amount of jitter to the predicted sensor curves, but we should also explore what happens when we increase the noise level. In general, a higher noise level *should* increase the difference between the unclipped error and clipped error, as more values are pushed out of bounds. In addition, it should make the resultant sensor curves even more jittery and obliterate more and more structure from those sensor curves until it is overcome completely by noise. Let's take a look at the graphs for noise with standard deviation 50 (both unclipped and clipped) and noise with standard deviation 100 (both unclipped and clipped), then at the chart of errors for all standard deviations in 10:10:100.




Figure 6: Actual and predicted sensors with noise of standard deviation 50, unclipped.




Figure 7: Actual and predicted sensors with noise of standard deviation 50, clipped between [0, 255].




Figure 8: Actual and predicted sensors with noise of standard deviation 100, unclipped.

Selected errors.:




Figure 9: Actual and predicted sensors with noise of standard deviation 100, clipped between $[0, 255]$.

	unclipped sensor RMSE	unclipped data RMSE	clipped sensor RMSE	clipped data RMSE
$\sigma = 0$	0.0	0.0000	0.0	0.0000
$\sigma = 1$	948.4	2.5089	949.1	2.5108
$\sigma = 2$	1887.3	4.9134	1877.8	4.8880
$\sigma = 3$	2930.7	7.7992	2873.6	7.6749
$\sigma = 4$	3596.6	9.5844	3425.1	9.3724
$\sigma = 5$	4685.1	12.4859	4227.3	11.6164
$\sigma = 6$	5473.3	14.5444	4810.9	13.6015
$\sigma = 7$	6728.6	18.0828	5574.2	17.6235
$\sigma = 8$	7566.5	19.7145	6126.3	18.8280
$\sigma = 9$	7925.0	21.6100	6177.6	21.9431
$\sigma = 10$	8797.3	23.4598	7276.3	25.0892

5.1 Discussion

The error appears to increase linearly with sigma. The plots also show that sigma causes the predicted sensors to be much more jittery. The curves even start to be harder to pick out among the jitter. As expected, more noise does increase the chance of clipping. However, clipping has a strange effect on the respective errors. While clipping increases the data error, it decreases the sensor error. I am a little stumped at this result, and I have double-checked my code to verify that it's doing what I think it's doing. If the results are correct, we could hypothesize as to why clipping would reduce the sensor error. One way to think of it might be to think in terms of our original data range and how we scaled it. At the beginning of the program, we scaled our random spectra values by a constant k at the beginning of the program to ensure that the RGB outputs would always be in range $[0, 255]$. Returning to our noisy samples, if we clip a negative number to 0, it *must* be closer to its original value because of that scaling constant k . I'm still stumped as to why the data error is larger, though.

6 (#5) Gamma Correction

Here, we simply solve the gamma correction equation for γ :

$$f(x) = 255\left(\frac{x}{255}\right)^{\frac{1}{\gamma}}$$

Substituting for x , and $f(x)$:

$$\begin{aligned}
 127 &= 255\left(\frac{80}{255}\right)^{\frac{1}{\gamma}} \\
 \frac{127}{255} &= \frac{80}{255}^{\frac{1}{\gamma}} \\
 \log\left(\frac{127}{255}\right) &= \frac{1}{\gamma}\log\left(\frac{80}{255}\right)
 \end{aligned}$$

$$\gamma = \frac{\log(\frac{80}{255})}{\log(\frac{127}{255})} = 1.6630$$

7 (#6) Problems Using the Pseudoinverse on Real Data




Figure 10: This is my best guess of an image of the given real-world RGB values. I wanted to visualize it a little better, so I guessed at the height and width.

When using this technique on real world data, we find that it performs very poorly.




Figure 11: Extremely poor performance of predicted sensor curves. The noise is so bad that the actual sensor curves are dwarfed and appear close to 0.

The errors are extremely high as well:

Sensor RMSE	Data RMSE
1.0138e6	5.7426

7.1 Discussion

Consider the case where none of the spectra samples have a non-zero value in a particular spectral bin. In this case, the matrix might as well be 100 columns instead of 101! Thinking in terms of linear algebra,

if we find that our spectral samples lie on some subspace of the actual space of possible spectral samples, we have a degenerative matrix. The basis of such a matrix is less than the dimensionality. For a normal n -by- n matrix, this would mean that the matrix can't be inverted. In the case of our pseudo-inverse, maybe something similar happens. While we can still take the pseudo-inverse, perhaps it's no longer useful.

8 (#7) Constrained Least Squared Error

If we constrain sensors to only have positive responses, the prospect is much less bleak. Here, we must use a different technique, adding a constraint expression (using `quadprog` in MATLAB). The resultant sensors are still very spikey, but the spikes tend to be near the actual sensor peaks.




Figure 12: Predicted sensor curves when we constrain them to be 0 or positive. The results are much better, and the spikes tend to be near the actual sensor peaks. However, there are still several outlier spikes that are nowhere near the peaks.

9 (#8) Smooth, Non-negative Predictions

We can also attempt to make the curves smoother to get rid of the spikes above. We can do so rather cleverly by concatenating the following matrix at the bottom of our spectra matrix:

$$M = \begin{bmatrix} 1 & -1 & 0 & \cdots & 0 & 0 \\ 0 & 1 & -1 & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \ddots & 0 \\ 0 & \cdots & \cdots & \cdots & 1 & -1 \end{bmatrix}$$

Now, this essentially represents the difference between adjacent spectral response buckets for our sensors. We want each of these differences to be small, so we append 0's at the bottom of the RGB output matrix. Finally, we multiply M by a constant λ to allow us to control the smoothness. Several different plots are presented for different lambda values. I think a value of $\lambda = 0.01$ matches the original sensors pretty well.




Figure 13: lambda (smoothing constant) = 0.0001 (very low)




Figure 14: λ (smoothing constant) = 0.001 (low)




Figure 15: λ (smoothing constant) = 0.01 (about right)




Figure 16: λ (smoothing constant) = 0.1 (high)




Figure 17: λ (smoothing constant) = 1 (very high)

9.1 Discussion

We can turn the λ knob so that the curves are as smooth or jagged as we want, but, in a real world case, we might not have the training data (actual sensor curves) to compare against. In that case, how would we

know a good value for λ ?