

CSC 577 HW2

Simon Swenson

1/23/2019

1 Introduction

This assignment was divided into two main sections, one focusing on model fitting and one focusing on reasoning about perspective in images. In the model fitting section, the focus was a comparison between non-homogeneous linear least squares fitting, using the now-familiar "dagger" of the input matrix, which we also explored in assignment 2, and a new method which treats the y coordinates of the input points equally to the x coordinates, which the non-homogeneous method did not use. In the image perspective system, we explore drawing parallel lines and making a judgment based on rules, namely, (1) that parallel lines in 3-d space lead to the same vanishing point in camera projection space and (2) that coplanar lines lead to the same horizon line. In the writeup, I use Θ to stand for the model parameters to prevent confusion with X (input points). As always, latex tends to float figures around in strange ways, so please check adjacent pages whenever figures are referenced or where you expect a figure to appear. All questions for the assignment were completed.

2 A Comparison of Line Fitting Strategies

We are given an arbitrary set of 2-d points and asked to find a good line fit to that set. However, there are many different ways to fit a line. For this assignment, we just explored two: non-homogeneous least squares and homogeneous least squares. The main difference between the two formulations is that the first formulation treats the y axis (and the y coordinates of the input points) as completely different from the x axis (and x coordinates). In this formulation, error is only minimized vertically from each point to the regression line. In the second, error is minimized by the tangent of the regression line from each point to the regression line.

In the first formulation, we found that the least squares error was minimized by the following solution. (Note that I used Θ for model parameters.) Let $U = [X|1]$. Then the solution is:

$$\Theta = (U^T U)^{-1} U^T y$$

In the second formulation, we found that the least squares error was minimized by letting the model parameters (a, b) equal the elements of the eigenvector with the smallest eigenvalue of $U^T U$, where

$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ x_2 - \bar{x} & y_2 - \bar{y} \\ \vdots & \vdots \\ x_N - \bar{x} & y_N - \bar{y} \end{bmatrix}$$

Note that, in this formulation, the columns of U are interchangeable. Thus, both x and y are treated the same. This is the key of the homogeneous least squares solution.

To compare the two models, we must come up with some way to convert the model parameters of one to the model parameters of the other. In other words, we must be able to convert a, b to slope, intercept and vice versa. For the first, we start with the equation:

$$[X|Y] \begin{bmatrix} a \\ b \end{bmatrix} = \vec{0}$$

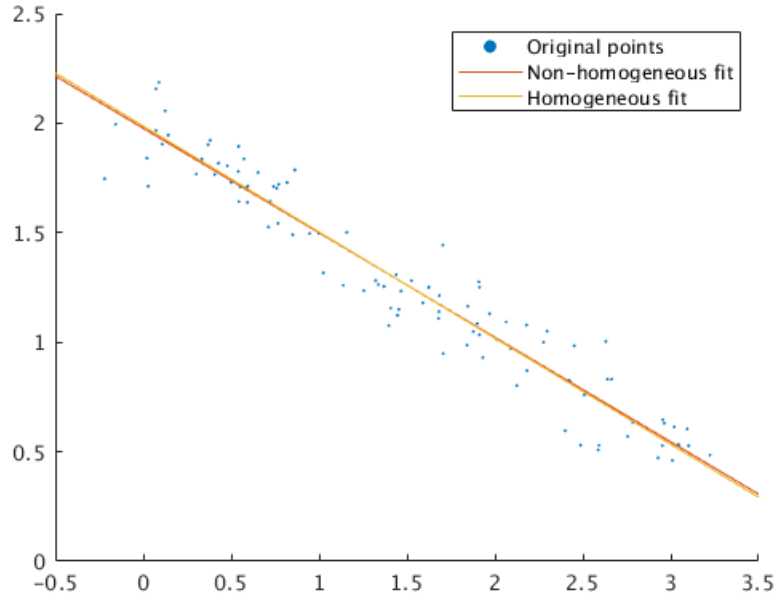


Figure 1: The original data points as a scatter plot, alongside both line fits. You can see from the figure that both line fits are very similar, but they are different, highlighting that each formulation leads to a different result.

Then we have:

$$\begin{aligned}
 & \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ x_2 - \bar{x} & y_2 - \bar{y} \\ \vdots & \vdots \\ x_N - \bar{x} & y_N - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \vec{0} \\
 & \begin{bmatrix} (x_1 - \bar{x})a + (y_1 - \bar{y})b \\ (x_2 - \bar{x})a + (y_2 - \bar{y})b \\ \vdots \\ (x_N - \bar{x})a + (y_N - \bar{y})b \end{bmatrix} = \vec{0} \\
 & a \begin{bmatrix} x_1 - \bar{x} \\ x_2 - \bar{x} \\ \vdots \\ x_N - \bar{x} \end{bmatrix} = -b \begin{bmatrix} y_1 - \bar{y} \\ y_2 - \bar{y} \\ \vdots \\ y_N - \bar{y} \end{bmatrix} \\
 & -\frac{a}{b} \begin{bmatrix} x_1 - \bar{x} \\ x_2 - \bar{x} \\ \vdots \\ x_N - \bar{x} \end{bmatrix} = \begin{bmatrix} y_1 - \bar{y} \\ y_2 - \bar{y} \\ \vdots \\ y_N - \bar{y} \end{bmatrix}
 \end{aligned}$$

$$\begin{aligned}
-\frac{a}{b} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix} - \bar{x} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} - \bar{y} \\
-\frac{a}{b} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix} + \frac{\bar{x}a}{b} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} - \bar{y} \\
-\frac{a}{b} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix} + \frac{\bar{x}a}{b} + \bar{y} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \\
\begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_N & 1 \end{bmatrix} \begin{bmatrix} -\frac{a}{b} \\ \frac{\bar{x}a}{b} + \bar{y} \end{bmatrix} &= \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}
\end{aligned}$$

In this form, it is clear that $m = -\frac{a}{b}$ and $b = \frac{\bar{x}a}{b} + \bar{y}$. Now, to convert in the other direction, we use the following equations, solving for a, b:

$$\begin{aligned}
m &= -\frac{a}{b} \\
a^2 + b^2 &= 1
\end{aligned}$$

First, we solve for a in the first equation:

$$a = -mb$$

Then, we substitute that a into the second equation:

$$\begin{aligned}
(-mb)^2 + b^2 &= 1 \\
m^2b^2 + b^2 &= 1 \\
b^2(m^2 + 1) &= 1 \\
b^2 &= \frac{1}{m^2 + 1} \\
b &= \pm \sqrt{\frac{1}{m^2 + 1}}
\end{aligned}$$

Finally, we plug that b into the first equation:

$$a = -m \pm \sqrt{\frac{1}{m^2 + 1}}$$

Now that we have a method to convert each model into the other, we can gather some results:

	Non-homogeneous Model	Homogeneous Model
Slope	-0.4772	-0.4840
Y-intercept	1.9769	1.9868
A	0.4307	-0.4357
B	0.9025	-0.9001
Non-homogeneous RMSE	0.1250	0.1252
Homogeneous RMSE	0.1128	0.1127

We see from the results that both models have a higher non-homogeneous RMSE than homogeneous RMSE. The following figure attempts to explain why that is the case:




Figure 2: A diagram illustrating non-homogeneous error (vertical distance) and homogeneous error (tangent distance). Notice that the homogeneous error is less. (Compare the circle radii.) In fact, this is the case for all fits, save for a horizontal line.

Now that we have explored the relationship between the two error metrics, let's explore the relationship between each error metric and both models. As we discussed in class, the non-homogeneous solution minimizes the mean squared error. Thus, we would expect any other solution, even the homogeneous solution, to have a higher non-homogeneous RMSE. This is indeed the case. A similar situation arises for homogeneous RMSE. We know by derivation that the homogeneous model will minimize the homogeneous mean squared error, and we can indeed see that that is the case based on our results.

3 Line Fitting - Graduate Discussion

Our observations above can be formalized below.

3.1 Theorem: Homogeneous MSE will always be less than or equal to non-homogeneous MSE

For a geometric proof sketch, refer to the nearby diagram (showing circles representing absolute distance). Since squared distance as a function of absolute distance is monotonically increasing (consider the right-hand side of a parabola), the diagram holds for squared distance as well as absolute distance. Thus, it applies to mean squared error. The diagram also holds for any angle of line and any data point offset. (When the line is horizontal, homogeneous MSE equals non-homogeneous MSE, but because of how we've stated the theorem, the theorem still holds.)

3.2 Theorem: The non-homogeneous model given above ensures a non-homogeneous error less than or equal to any possible linear model

When discussing minimum values, it is always a good idea to turn to the idea of derivatives. Doing derivatives in vector space is a bit tricky, however. We turn to the idea of the gradient, which uses derivatives in the form of a vector. The gradient is a vector that always points in the direction of greatest positive change. Note that, at a peak (maximum) or valley (minimum), there is no direction of greatest positive change, so the gradient is 0. Recall that the gradient is defined as:

$$\text{grad}(f(m, b)) = \begin{bmatrix} \frac{\delta f}{\delta m} \\ \frac{\delta f}{\delta b} \end{bmatrix}$$

Note that we treat m and b as unknowns in this formulation: We flip the roles of the data-points (X, y) and the model parameters when we "train." (Although I am being a bit loose with the word "train" in this scenario.)

Let us return to the original error formulation and apply the gradient in terms of m and b:

$$\text{homogeneous_mse} = \frac{\sum_{1..N} (x_i m + b - y_i)^2}{N}$$

Then:

$$\begin{aligned} \text{grad}(\text{homogeneous_mse}) &= \begin{bmatrix} \frac{\delta \text{homogeneous_mse}}{\delta m} \\ \frac{\delta \text{homogeneous_mse}}{\delta b} \end{bmatrix} \\ &= \begin{bmatrix} \frac{\delta}{\delta m} \left(\frac{\sum_{1..N} (x_i m + b - y_i)^2}{N} \right) \\ \frac{\delta}{\delta b} \left(\frac{\sum_{1..N} (x_i m + b - y_i)^2}{N} \right) \end{bmatrix} \end{aligned}$$

Let's first consider the derivative wrt m:

$$\frac{\delta \text{homogeneous_mse}}{\delta m} = \frac{2}{N} \sum_{1..N} ((x_i m + b - y_i) x_i)$$

Similarly, we figure the derivative wrt b:

$$\frac{\delta \text{homogeneous_mse}}{\delta b} = \frac{2}{N} \sum_{1..N} (x_i m + b - y_i)$$

To find the extrema, we set the gradient to the zero vector and solve:

$$\begin{aligned} \vec{0} &= \frac{2}{N} \begin{bmatrix} \sum_{1..N} ((x_i m + b - y_i) x_i) \\ \sum_{1..N} (x_i m + b - y_i) \end{bmatrix} \\ \vec{0} &= \frac{2}{N} \sum_{1..N} \begin{bmatrix} (x_i m + b - y_i) x_i \\ x_i m + b - y_i \end{bmatrix} \\ \vec{0} &= \frac{2}{N} \sum_{1..N} \begin{bmatrix} x_i^2 m + x_i b - x_i y_i \\ x_i m + b - y_i \end{bmatrix} \\ \vec{0} &= \sum_{1..N} \begin{bmatrix} x_i^2 m + x_i b \\ x_i m + b \end{bmatrix} - \sum_{1..N} \begin{bmatrix} x_i y_i \\ y_i \end{bmatrix} \\ \sum_{1..N} \begin{bmatrix} x_i y_i \\ y_i \end{bmatrix} &= \sum_{1..N} \begin{bmatrix} x_i^2 m + x_i b \\ x_i m + b \end{bmatrix} \\ \sum_{1..N} \begin{bmatrix} x_i y_i \\ y_i \end{bmatrix} &= \sum_{1..N} \begin{bmatrix} x_i^2 & x_i \\ x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \\ \begin{bmatrix} \sum_{1..N} x_i y_i \\ \sum_{1..N} y_i \end{bmatrix} &= \begin{bmatrix} \sum_{1..N} x_i^2 & \sum_{1..N} x_i \\ \sum_{1..N} x_i & N \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \end{aligned}$$

From here, we use standard techniques to find the inverse matrix of the matrix on the right (which we are multiplying by (m, b) on the left). As the expression gets quite unmanageable otherwise, I will omit the subscripts for \sum, x, y . They are assumed.

$$\frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \begin{bmatrix} \sum xy \\ \sum y \end{bmatrix} = \frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \begin{bmatrix} \sum x^2 & \sum x \\ \sum x & N \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix}$$

$$\frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \begin{bmatrix} \sum xy \\ \sum y \end{bmatrix} = \begin{bmatrix} m \\ b \end{bmatrix}$$

Now, we will attempt to derive the same expression from the formulation of the non-homogeneous solution presented in class:

$$\begin{bmatrix} m \\ b \end{bmatrix} = (U^T U)^{-1} U^T Y$$

We begin:

$$\begin{aligned} U &= \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_N & 1 \end{bmatrix} \\ U^T U &= \begin{bmatrix} x_1 & x_2 & \cdots & x_N \\ 1 & 1 & \cdots & 1 \end{bmatrix} \begin{bmatrix} x_1 & 1 \\ x_2 & 1 \\ \vdots & \vdots \\ x_N & 1 \end{bmatrix} \\ &= \begin{bmatrix} \sum x^2 & \sum x \\ \sum x & N \end{bmatrix} \\ (U^T U)^{-1} &= \frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \\ (U^T U)^{-1} U^T Y &= \frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \begin{bmatrix} x_1 & x_2 & \cdots & x_N \\ 1 & 1 & \cdots & 1 \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \\ &= \frac{1}{N \sum x^2 - \sum x \sum x} \begin{bmatrix} N & -\sum x \\ -\sum x & \sum x^2 \end{bmatrix} \begin{bmatrix} \sum xy \\ \sum y \end{bmatrix} \end{aligned}$$

The two expressions are the same, so it is thus proven.

Note: To complete the proof, we would need to verify that we have actually found a minimum based on concavity, but that's a bit tedious for this assignment, so I will leave it as is.

3.3 Theorem: The homogeneous model given above ensures a homogeneous error less than or equal to any possible linear model

The proof sketch follows from the optional material found in slide deck 5. We initially want to find x in the following:

$$Ux \approx \vec{0}$$

with the constraint:

$$\|x\| = 1$$

Note that the equation above has no exact solution, so we want to minimize Ux . Since Ux is not square, we multiply on the left by the transpose, similar to how we handled the non-homogeneous solution. This expression is the squared error. Then we want to minimize:

$$\min(x^T U^T U x)$$

with the constraint:

$$x^T x = 1$$

We can rewrite $U^T U$ based on its eigen decomposition. Then the expression above becomes:

$$\min(x^T V \Lambda V^T x)$$

It can be shown that $U^T U$ is symmetric. Furthermore, it can be shown that eigenvectors of symmetric matrices are orthogonal. Then the dot product between of those distinct eigenvectors is 0. If we let x equal the eigenvector that corresponds to the smallest eigenvalue, $V^T x$ becomes a vector with one 1 and the remaining entries 0 because one entry is the dot product of a unit vector with itself and all other entries are the dot product between orthogonal vectors and are thus 0. Note that this is the case *for any eigenvector that we choose*. A similar vector arises on the other side ($x^T V$). Thus, in this formulation, the only real thing we have control over is selecting which entry in Λ that we select (by picking a different eigenvector). Obviously, we want the smallest, so we pick the eigenvector that corresponds to the smallest eigenvalue.

This is an argument for why, if we restrict our search to only eigenvectors, we should pick the eigenvector with the smallest eigenvalue, but why restrict our picks to the eigenvectors, in the first place? Why couldn't the solution be some other random unit vector? Well, it has to do with the way eigenvectors work. Building off of what we learned about PCA from a previous assignment, the eigenvector with the largest eigenvalue will point in the direction of maximum variance. For a roughly linear distribution, this is the same direction as the line, itself. However, we learned from class that a, b is orthogonal to the line that we want. Knowing that eigenvectors of symmetric matrices are orthogonal, by process of elimination, we must pick the eigenvector with the smallest eigenvalue to be a, b .

3.4 Discussion: What if our solution is just a local minimum?

This is a familiar concept to anyone familiar with machine learning. Often, the error function does not have a single extremum. If our solution only gives a local minimum, we have no good way of knowing if that minimum is also a global minimum, assuming all minima cannot be solved for explicitly. (If we have the set of all minima, we can do calculations to find out which is the global minimum.) Often, the set of all minima cannot be solved for explicitly, in which case gradient descent over the error function is often used. Starting from an arbitrary point, we repeatedly walk in the opposite direction of the gradient of the error function until we reach a minima. Intuitively, since we are taking steps, the process only knows about the immediate surroundings of our current point. It doesn't know the layout of the entire error function. Thus, it is impossible to know if we've reached a global minimum after such a process.

4 Reasoning About Perspective: Skyscraper

This image is fairly easy to reason about because the skyscraper is a fairly regular object. Here, we can easily draw the first two sets of parallel lines: vertical along the front face and horizontal along the front face. The third set of points can be generated by counting windows (six windows higher on the left-hand side than the right-hand side). All sets of lines do roughly meet at corresponding common points. In addition, we find that those points are all co-linear. Thus, we conclude that the image is in proper perspective.

5 Reasoning About Perspective: Chandelier

The candels of this chandelier look to be aligned in a hexagon. If we assume that the hexagon is fairly regular, we can assume that the opposite edges are parallel in world coordinates. When we draw those corresponding lines, most do, indeed, meet at a point in projection space. However, one set does not. In addition, when we do this for all three pairs of parallel lines of the hexagon, we see that the three corresponding intersection points are not co-linear in projection space, which we would expect, since these points are all co-planar in world space. Thus, we can conclude that the image is *not* in proper perspective.




Figure 3: A skyscraper. By following parallel, coplanar lines in world space, we find that parallel lines meet at common points and sets of coplanar, parallel lines produce common points along a common horizon line.




Figure 4: A painting(?) of a chandelier. By following parallel, coplanar lines in world space, we find that the resultant rays in projection space are not co-linear.