

# CSC 577 HW7

Simon Swenson

3/20/2019

## 1 Introduction




Figure 1: The four scenes explored in this assignment, lit by a white light, gamma corrected, and exposure increased.

In this assignment, we were presented various scenes lit by two different lights: a deep blue light and a pure white light. We explored various ways to estimate the red, green, and blue values for these lights, including using white patches, using the "max RGB" approach, using the "gray world" approach, and minimizing the mean squared error, and using a built-in MATLAB solver to find the minimum solution to RMSE r,g. Each of these methods has advantages and disadvantages, which we saw on visual inspection, by calculating the angular error, and by calculating the RMSE r,g.

## 2 Estimating Light Values from White Patches

Similar to how we used a regular grid over  $R^3$  to calibrate a camera, we can use a plain white patch to find the red, green, and blue values for a light. This is essentially an "oracle" approach, since we manually select the pixels which are within the white patch and know that that patch is pure white. This yields the following two colors for the lights:

	Red	Green	Blue
syl-50MR16Q (white light)	239	222	250
solux-4100 (blue light)	131	159	250

The angular error between these two light colors is 0.2416. From there, we use the diagonal color model to transform each pixel in the source image, using the following equation for the diagonal color model:

$$\begin{bmatrix} R_2 \\ G_2 \\ B_2 \end{bmatrix} = \begin{bmatrix} \frac{R_{L2}}{R_{L1}} & & \\ & \frac{G_{L2}}{G_{L1}} & \\ & & \frac{B_{L2}}{B_{L1}} \end{bmatrix} \begin{bmatrix} R_1 \\ G_1 \\ B_1 \end{bmatrix}$$

Using this approach yields good results, but the corrected image appears slightly darker than the target image. However, upon closer inspection, the white patch has the correct average color. The RMSE r,g results between the original two images and the predicted image and the target image are shown below. The prediction image decreases the RMSE r,g by more than half:

blue-to-white RMSE r,g	0.0746
predicted-to-white RMSE r,g	0.0335

However, qualitatively, even though the RMSE r,g was only reduced by a little more than half, the image looks *much* closer to the target image.



Figure 2: Left: A color swatch lit by the blue light. This is the source. Center: The color swatch transformed from the blue light space to the white light space using the diagonal color model. This is the prediction. Right: A color swatch lit by the white light. This is the target.

What could be the reason for the white patch being correct, but the rest of the scene appearing darker? One possible explanation is that the patch is not pure white. Any impurity of the color could cause the transformation to fail, since it will not give an accurate reading for the light values.

### 3 Estimating Light Values using Max RGB

Now that we have a good baseline for how well we could calculate the light values in an oracle situation, we can compare two different approaches to that one and see how well we can do without perfect information. The first approach is called "max RGB." The algorithm finds the maximum value in each of the three channels of the image and uses that as the value for the light:

$$\begin{bmatrix} R_W^U \\ G_W^U \\ B_W^U \end{bmatrix} = \begin{bmatrix} \max_{pix} R \\ \max_{pix} G \\ \max_{pix} B \end{bmatrix}$$

Intuitively, this approach should fair well when the image contains dielectric materials, since such materials have specular highlights, which should match the color of the light closely. This should also work well when there is any sort of white patch in the image, since that will dominate the "max" values. We explore this approach using three different scenes, a scene with several apples and a blue background, a scene with a ball and a black background, and a scene with blocks and a red background. Doing so leads to the following predicted lights:

	Red	Green	Blue
Apples Light	182	228	250
Ball Light	112	143	250
Blocks Light	250	242	209
and the quantitative results:			
	Angular Error	RMSE r,g Error	
Apples Light	0.1713	0.0431	
Ball Light	0.0637	0.0661	
Blocks Light	0.3484	0.0746	

And the qualitative results:




Figure 3: Left: The three scenes (apples, ball, blocks) lit by the blue light. This is the source. Center: The scenes transformed from the blue light space to the white light space using the max RGB light value and the diagonal color model. This is the prediction. Right: The scenes lit by the white light. This is the target.

This approach did very well in the case of the apples. Since the surface of the apples is shiny, that shean provides a good value for the light. Unexpectedly, though, the approach did not do as well on the ball image, even though the ball has a large patch of white in it. Inspecting the image closer reveals the culprit, though: a vertical, bright aberration to the bottom-left of the ball. I can only guess that this has affected the result, since it is brighter than the white on the ball, itself. Lastly, the approach didn't seem to change the blocks image much. It seems to have increased the amount of blue in the image slightly, since the red now looks more purplish, the yellow looks more washed out, and the blue looks a bit brighter. In other words, the max RGB algorithm found that the light was close to gray, but slightly yellowish, since it made the image more blue, and blue is the opposite of yellow. This is best explained by the composition of the image, which contains very bright yellow blocks, but the other blocks are fairly muted in color, in comparison.

## 4 Estimating Light Values using Gray World

The second possible approach is the gray world approach. This approach takes the average value of each color channel rather than the max value:

$$\begin{bmatrix} R_W^U \\ G_W^U \\ B_W^U \end{bmatrix} = \begin{bmatrix} avg_{pix}R \\ avg_{pix}G \\ avg_{pix}B \end{bmatrix}$$

Intuitively, we would expect a scene that's gray to produce good results with this algorithm. It may suffer when large portions of the image are dominated by a non-gray object (which we'll see later).

The lights recovered using this method follow:

	Red	Green	Blue
Apples Light	50	79	250
Ball Light	125	156	250
Blocks Light	250	155	179

and the quantitative results:

	Angular Error	RMSE r,g Error
Apples Light	0.3422	0.1021
Ball Light	0.0244	0.0635
Blocks Light	0.4091	0.1159

And the qualitative results:




Figure 4: Left: The three scenes (apples, ball, blocks) lit by the blue light. This is the source. Center: The scenes transformed from the blue light space to the white light space using the gray world light value and the diagonal color model. This is the prediction. Right: The scenes lit by the white light. This is the target.

The results using this method are pretty bad for the apples scene and the blocks scene. (However, it seems to do alright with the ball scene.) An understanding of how gray world works can tell us why the results are so bad and what's going on. Since gray world averages over all of the pixel values, larger color patches will tend to dominate the calculated value. In the first image, the background image is blue as lit by

the white light. This brings the calculated light to be much bluer than it should be. (If the background were roughly gray, things would be alright.) To counteract this, the predicted result adds way too much yellow light (the opposite of blue). This makes the apples appear too red.

The blocks image has a similar problem because the background is red. The algorithm finds that the light color is actually red, due to this, and to counteract the red light, the prediction has way too much cyan added (the opposite of red).

The ball is unaffected by this, since its background is roughly gray/black. The results for the ball are actually quite good, with a better angular error than max RGB and a RMSE r,g value on par with max RGB.

## 5 Minimizing Mean Squared Error using the Gradient

The holy grail of many types of problems is the minimization of an error metric. Here, we consider the MSE metric:

$$E = \frac{\sum_y^H \sum_x^W \sum_{c \in \{R, G, B\}} (\frac{L_c^T}{L_c^S} S_{x,y,c} - T_{x,y,c})^2}{3HW}$$

Where  $H$  is the image height,  $W$  is the image width,  $L^T$  is the target light,  $L^S$  is the source light,  $S_{x,y,c}$  is the value of the source image at pixel  $x, y$  for channel  $c$ , and  $T_{x,y,c}$  is the value of the target image at pixel  $x, y$  for channel  $c$ .

Our objective is to find the source light values that minimize this function:

$$\text{argmin}_{L_R^S, L_G^S, L_B^S} \left( \frac{\sum_y^H \sum_x^W \sum_{c \in \{R, G, B\}} (\frac{L_c^T}{L_c^S} S_{x,y,c} - T_{x,y,c})^2}{3HW} \right)$$

To do so, we compute the gradient of that equation with respect to the three source light channels, then set that gradient to the 0 vector, then solve. (Note that I think the non-homogeneous LSE method using the pseudoinverse would work as well, but I used the gradient method instead. I think the two answers are equivalent.) Each channel is symmetric and will have a very similar partial derivative, so let's just consider  $\frac{\delta E}{\delta L_R^S}$ :

$$\begin{aligned} \frac{\delta E}{\delta L_R^S} &\propto - \sum_y^H \sum_x^W \left( \frac{L_R^T}{L_R^S} S_{x,y,R} - T_{x,y,R} \right) \left( \frac{L_R^T}{(L_R^S)^2} S_{x,y,R} \right) \\ &\propto \sum_y^H \sum_x^W \left( \frac{L_R^T}{L_R^S} (S_{x,y,R})^2 \right) - \sum_y^H \sum_x^W (T_{x,y,R} S_{x,y,R}) \end{aligned}$$

Setting the gradient to 0, we have:

$$\begin{aligned} \sum_y^H \sum_x^W (T_{x,y,R} S_{x,y,R}) &= \sum_y^H \sum_x^W \left( \frac{L_R^T}{L_R^S} (S_{x,y,R})^2 \right) \\ \sum_y^H \sum_x^W (T_{x,y,R} S_{x,y,R}) &= \frac{L_R^T}{L_R^S} \sum_y^H \sum_x^W (S_{x,y,R})^2 \\ L_R^S &= L_R^T \frac{\sum_y^H \sum_x^W (S_{x,y,R})^2}{\sum_y^H \sum_x^W (T_{x,y,R} S_{x,y,R})} \end{aligned}$$

Similar for the remaining two channels. This is the expression which we will solve for. At this point, it should be noted that, since all three channels are independent of each other, and each individual result in the sum is squared and thus is guaranteed to be positive, surely minimizing the error result for all three

channels will also minimize the error result in just the r,g channels, so this *should* also minimize the RMSE r,g metric. Computing this minimum gives us the following results:

The lights recovered using this method follow:

	Red	Green	Blue
Apples Light	131	159	250
Ball Light	133	163	250
Blocks Light	126	157	250

and the quantitative results:

	Angular Error	RMSE r,g Error
Apples Light	0.0025	0.0159
Ball Light	0.0084	0.0631
Blocks Light	0.0158	0.0230

And the qualitative results:




Figure 5: Left: The three scenes (apples, ball, blocks) lit by the blue light. This is the source. Center: The scenes transformed from the blue light space to the white light space using the minimum MSE light value and the diagonal color model. This is the prediction. Right: The scenes lit by the white light. This is the target.

The predictions are remarkably good. The ball result is a little worse than expected, due to the aberration to the bottom-left of the ball. Examining the light values show that the predicted lights are very close to each other for each scene. The error metrics are also very good.

## 6 Off-the-shelf Optimization Method

Matlab provides a function for finding minimum values for a given, arbitrary function. One such method is "simulannealbnd," which implements the simulated annealing strategy, and enforces a given bound on the final output vector. The simulated annealing strategy is good for finding a global minimum, as opposed to a local minimum. We can simply use the RMSE r,g function that we already have defined to do so. Comparing

the results should give us some more evidence for whether minimizing the MSE actually minimizes the RMSE r,g as well.

The lights recovered using this method follow:

	Red	Green	Blue
Apples Light	128	156	250
Ball Light	146	172	250
Blocks Light	135	159	250

and the quantitative results:

	Angular Error	RMSE r,g Error
Apples Light	0.0130	0.0141
Ball Light	0.0434	0.0606
Blocks Light	0.0060	0.0148

The results are a bit unexpected and reveal that the minimum MSE measurement does not actually minimize the RMSE r,g values, since these are slightly better. I had originally thought that, since the three color channels are independent, the answer which minimizes the MSE would also minimize the RMSE r,g value. However, taking a closer look at the RMSE r,g metric reveals that the colors are *not* actually independent in this metric:

$$r = \frac{R}{R + G + B}$$

$$g = \frac{G}{R + G + B}$$

$$b = \frac{B}{R + G + B}$$

Since this metric essentially takes proportion, rather than raw value, into account, the channels are not independent, so the answer which minimizes the MSE will be different from the answer that minimizes the RMSE r,g.